

Sleep Stage Classification from Wrist ACC + Heart Rate (PhysioNet Sleep-Accel / Apple Watch)

Chouaib (github.com/ChouaibB)

January 27, 2026

Contents

1. Data source: Sleep-Accel (PhysioNet; Apple Watch + PSG)	2
2. Data loading and check	3
3. Epoch alignment (30s) — assign raw samples to PSG epochs	4
4. Feature extraction	6
4.1 Minimal per-epoch feature summaries (ACC magnitude + HR)	6
4.2 Per-subject missingness (% of epochs)	7
4.4 Causal history features (past-only rolling summaries)	7
4.5 Feature relevance and visualization	8
5. Modeling	10
6. Results & reporting	11
6.1 Performance summary	12
6.2 Confusion matrices	13

1. Data source: Sleep-Accel (PhysioNet; Apple Watch + PSG)

Dataset: *Motion and heart rate from a wrist-worn wearable and labeled sleep from polysomnography* (PhysioNet, v1.0.0)

- **PhysioNet dataset page (download + description):**
 - <https://physionet.org/content/sleep-accel/1.0.0/>
- **DOI (v1.0.0):**
 - <https://doi.org/10.13026/hmhs-py35>
- **Local path (expected):** download + unzip into `./data/sleep_accel/` (the `data/` directory is not committed to git)
 - expected: `heart_rate/`, `motion/`, `labels/`, `steps/`, plus `LICENSE.txt`
- **License (for files):** Open Data Commons Attribution License v1.0 (**ODC-By 1.0**)

Citations (as requested by PhysioNet):

- Walch, O. (2019). *Motion and heart rate from a wrist-worn wearable and labeled sleep from polysomnography* (version 1.0.0). PhysioNet. <https://doi.org/10.13026/hmhs-py35>
- Walch, O., Huang, Y., Forger, D., Goldstein, C. (2019). *Sleep stage prediction with raw acceleration and photoplethysmography heart rate data derived from a consumer wearable device*. SLEEP. <https://doi.org/10.1093/sleep/zsz180>
- Goldberger, A. L., et al. (2000). *PhysioBank, PhysioToolkit, and PhysioNet: Components of a new research resource for complex physiologic signals*. Circulation.

Files used (and expected columns)

We will use the following folders/files from the PhysioNet release:

- **Motion (ACC):** `motion/[subject]_acceleration.txt`
Columns per line: `t_sec`, `ax_g`, `ay_g`, `az_g`
where `t_sec` is seconds since PSG start, and accelerations are in `g`.
- **Heart rate (HR):** `heart_rate/[subject]_heartrate.txt`
Columns per line: `t_sec`, `hr_bpm`
where `hr_bpm` is heart rate in **beats per minute**.
- **PSG sleep labels:** `labels/[subject]_labeled_sleep.txt`
Columns per line: `t_sec`, `stage` with stage codes:
`Wake=0`, `N1=1`, `N2=2`, `N3=3`, `REM=5` (we drop unscored/invalid epochs if present).

Note: The dataset also includes `steps/` files, but we won't use them in the first version.

Notebook intention

Goal: build a **reproducible sleep-staging pipeline** from **wrist ACC + HR** aligned to **PSG-scored 30-second epochs**, with **leakage-aware, subject-wise evaluation**.

What we do:

1. Define the modeling unit as the **PSG epoch (30s)** and build one feature row per epoch.

2. Align wrist **ACC** and **HR** to each labeled 30s epoch (aggregate samples falling in $[t, t+30s)$), and attach the PSG stage label at t .
3. Extract simple, readable features per epoch:
 - ACC: magnitude and axis statistics + activity intensity proxies
 - HR: summary statistics + missingness indicators
4. Add causal context (history) features using past-only rolling summaries over recent epochs (e.g., last few minutes) to capture local sleep continuity without using future information.
5. Train and compare a small set of classical models using subject-wise cross-validation (GroupKFold) and report robust staging metrics (macro-F1, balanced accuracy, confusion matrices, per-subject performance).

Note on extra “pre-PSG” wearable data:

This dataset includes wearable streams that may start before **PSG time zero** (e.g., steps for days prior, HR for hours prior, motion shortly before). For the main staging pipeline we **restrict to the PSG-labeled interval** and only aggregate sensor data within labeled 30s epochs. Pre-PSG data can be used in extensions as **subject-level context** (e.g., prior-days activity summaries computed strictly from $t < 0$), but coverage varies across subjects and adds preprocessing complexity.

2. Data loading and check

We load per-subject **wrist accelerometer (ACC)**, **heart rate (HR)**, and **PSG sleep labels** from the PhysioNet Sleep-Accel folder structure.

Key conventions: - All timestamps are de-identified and expressed as **seconds since PSG start** ($t=0$ is the PSG start time). - For the main staging pipeline, we **restrict to the PSG-labeled interval** and only aggregate sensor samples that fall inside each labeled 30s epoch.

```
Found 31 subjects (from labels/). Example IDs: ['1066528', '1360686', '1449548', '1455390', '1818471']
```

Loaded ALL data:

```
labels: (27211, 3) | subjects=31
hr     : (254425, 3)    | subjects=31
acc    : (51819151, 5)   | subjects=31
```

Example subject 1066528:

```
labels_ex: (952, 3) | t range: [0.0, 28530.0]
hr_ex    : (16617, 3)    | t range: [-355241.7, 34491.2]
acc_ex   : (1281000, 5)   | t range: [-21684.8, 28626.5]
```

	stage_code	stage_name	count
0	0	Wake	185
1	1	N1	97
2	2	N2	299

3	3	N3	62
4	5	REM	309

3. Epoch alignment (30s) — assign raw samples to PSG epochs

We use PSG labels as the **epoch grid** (30s cadence), and assign raw HR and ACC samples to each epoch.

Output: one row per epoch with:

- epoch_id, t_start_sec, t_end_sec, stage_code
- raw per-epoch sequences: HR (t_sec, hr_bpm) and ACC (t_sec, ax_g, ay_g, az_g)

Note on sampling rates and interpolation:

ACC is high-rate, HR is sparse/irregular, and labels are per 30s epoch. We **do not resample or interpolate** signals to a common grid. Instead, we compute **per-epoch summary features** (e.g. counts, mean/std/percentiles, simple trends) directly from the raw samples assigned to each epoch at the next section.

epoch_raw: (26417, 11)

```

    subject_id  t_start_sec  stage_code  t_end_sec  epoch_id  \
0      1066528          0.0          0       30.0       0
1      3509524          0.0          0       30.0       1
2      4018081          0.0          0       30.0       2
3      4426783          0.0          0       30.0       3
4      5132496          0.0          0       30.0       4

                                         hr_t_sec  \
0  [6.38561010361, 6.38561010361, 6.38561010361, ...
1  [2.97270989418, 6.97270989418, 11.9727399349, ...
2  [1.20449995995, 5.2044699192, 9.2044699192, 18...
3  [3.28572010994, 7.28574991226, 13.2857499123, ...
4  [6.48358011246, 8.48358011246, 12.4835801125, ...

                                         hr_bpm  \
0  [52.0, 52.0, 52.0, 52.0, 52.0, 52.0, 51.0, 51...
1  [82.0, 84.0, 86.0, 86.0, 86.0, 83.0, 76.0]
2  [69.0, 66.0, 62.0, 57.0, 58.0, 61.0]
3  [71.0, 71.0, 72.0, 61.0, 60.0, 60.0]
4  [113.0, 115.0, 112.0, 117.0, 113.0]

                                         acc_t_sec  \
0  [0.0159480571747, 0.0360059738159, 0.055885076...
1  [0.00270414352417, 0.0179829597473, 0.03272795...
2  [0.0129339694977, 0.0286469459534, 0.042945861...
3  [0.0104320049286, 0.0303421020508, 0.050279140...
4  [0.0128321647644, 0.0269389152527, 0.042412996...

                                         acc_x  \
0  [0.4039307, 0.4039154, 0.4049072, 0.4083557, 0...

```

```

1 [0.0323334, 0.0274353, 0.02388, 0.0252228, 0.0...
2 [-0.4547424, -0.4547424, -0.455246, -0.4581757...
3 [-0.7835388, -0.7860413, -0.7800598, -0.783889...
4 [-0.1996918, -0.4198914, -0.3826447, -0.204345...

```

```

acc_y \
0 [0.4490051, 0.4480286, 0.4465485, 0.447525, 0...
1 [-0.4576111, -0.4629974, -0.4717712, -0.465408...
2 [0.5541077, 0.5541077, 0.5560608, 0.5536194, 0...
3 [-0.21698, -0.2355804, -0.2243347, -0.2145691,...
4 [0.3026581, 0.1375885, 0.0600433, 0.2941742, 0...

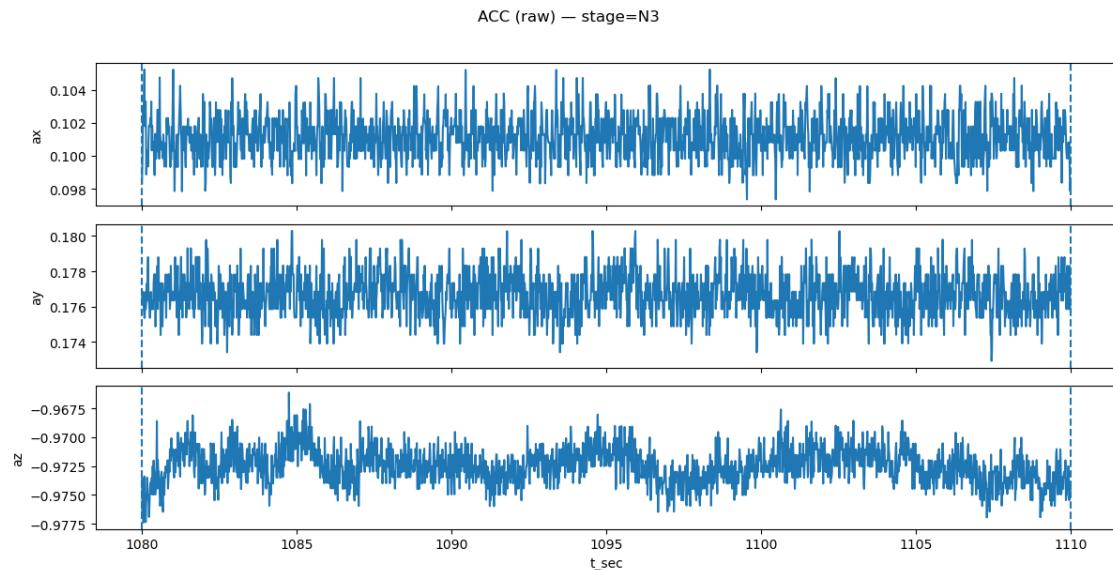
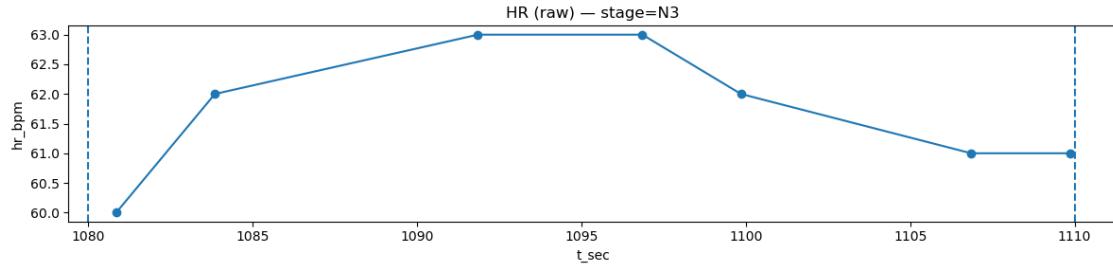
```

```

acc_z
0 [-0.7968597, -0.7953949, -0.7958527, -0.796768...
1 [-0.8885193, -0.8871613, -0.8794403, -0.874557...
2 [-0.6973724, -0.6978607, -0.6973877, -0.696945...
3 [-0.5306091, -0.5212555, -0.5314789, -0.546798...
4 [-0.9453125, -0.8695831, -0.9243622, -0.813476...

```

Example plot for one epoch: subject_id=759667 | epoch_id=756 | window=[1080.0,1110.0) | stage=N3



4. Feature extraction

We convert each PSG-aligned 30-second epoch into a single feature vector.
 To keep the notebook **simple and robust** (and to avoid a “feature zoo”, especially once we add **causal history features**), we use a **small, standard set of summary statistics**.

Signals and features

Accelerometer (ACC):

- We collapse tri-axial ACC into **magnitude** $|a| = \sqrt{ax^2 + ay^2 + az^2}$.
- Per epoch we compute: `count`, `mean`, `std`, `median`, `IQR`, `p10`, `p90`, `min`, `max` - Dynamic proxy `std(diff(|a|))` (often useful for “movement bursts”)

Heart rate (HR):

- Per epoch we compute: `count`, `mean`, `std`, `IQR`, `min`, `max`

Time

- `time_since_start_hours` (simple time context within the night)

Note on sampling and reliability ACC is high-rate while HR is sparse/irregular; labels are per 30s epoch. We **do not resample or interpolate**.

Summary statistics are computed directly from the raw samples assigned to each epoch, and `count` acts as a lightweight **signal reliability** indicator (few/no samples \rightarrow less reliable features).

4.1 Minimal per-epoch feature summaries (ACC magnitude + HR)

`epoch_features` shape: (26417, 22)

	subject_id	epoch_id	t_start_sec	t_end_sec	stage_code	\
0	1066528	0	0.0	30.0	0	
1	3509524	1	0.0	30.0	0	
2	4018081	2	0.0	30.0	0	
3	4426783	3	0.0	30.0	0	
4	5132496	4	0.0	30.0	0	
	time_since_start_hours		hr_mean	hr_std	hr_iqr	hr_min ... \
0	0.0	52.166667	0.897527	0.00	51.0	...
1	0.0	83.285714	3.325842	3.50	76.0	...
2	0.0	62.166667	4.219663	6.25	57.0	...
3	0.0	65.833333	5.520165	10.75	60.0	...
4	0.0	114.000000	1.788854	2.00	112.0	...
	acc_mag_count	acc_mag_mean	acc_mag_std	acc_mag_median	acc_mag_iqr	\
0	1500	1.000029	0.001947	1.000037	0.002023	
1	1848	0.993591	0.034502	0.992176	0.010523	
2	1918	0.999597	0.001879	0.999510	0.002095	
3	1716	1.000140	0.081830	0.991798	0.012816	
4	1861	1.019645	0.174913	0.998606	0.020979	

	acc_mag_p10	acc_mag_p90	acc_mag_min	acc_mag_max	acc_mag_diff_std
0	0.997970	1.002068	0.977874	1.017645	0.002779
1	0.967860	1.019650	0.780282	1.235724	0.019853
2	0.997544	1.001872	0.990996	1.008555	0.001812
3	0.972162	1.025978	0.512724	2.283005	0.063535
4	0.932608	1.136894	0.288624	3.711552	0.122539

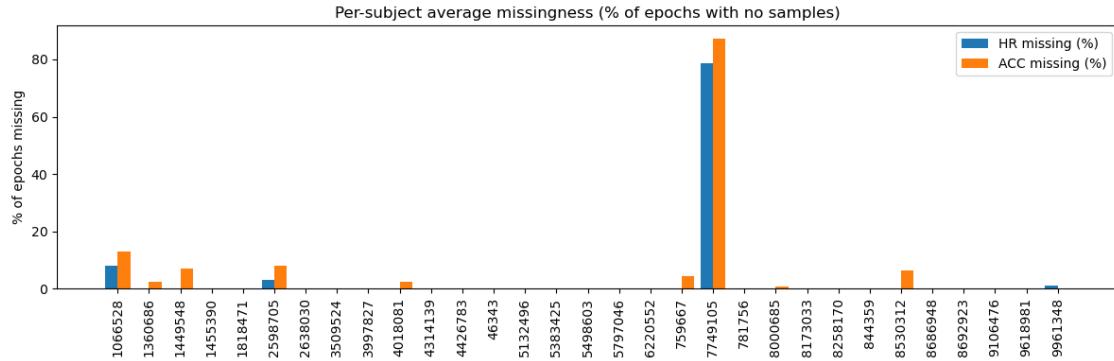
[5 rows x 22 columns]

4.2 Per-subject missingness (% of epochs)

We quantify how often each signal is missing at the epoch level:

- **HR missing:** `hr_count == 0`
- **ACC missing:** `acc_mag_count == 0`

We plot the **percentage of missing epochs per subject** for HR and ACC to spot subjects with low signal coverage.



4.4 Causal history features (past-only rolling summaries)

Sleep is a continuous process, so we add lightweight **history context** to each epoch using only **past information**.

For each subject, we compute history (rolling means) for the main **level** and **variability** signals `hr_mean`, `hr_std`, `acc_mag_mean`, `acc_mag_std`, over the previous:

- **5 epochs** (~2.5 minutes)
- **20 epochs** (~10 minutes)

To keep this strictly causal, we use `shift(1)` so the rolling window for epoch t uses epochs $< t$ only.

```
Added history features: ['hr_mean_roll15', 'hr_std_roll15', 'acc_mag_mean_roll15',
'acc_mag_std_roll15', 'hr_mean_roll20', 'hr_std_roll20', 'acc_mag_mean_roll20',
'acc_mag_std_roll20']
epoch_features shape with history features: (26417, 30)
```

```

subject_id epoch_id t_start_sec t_end_sec stage_code \
0 1066528 0 0.0 30.0 0
1 1066528 11 30.0 60.0 0
2 1066528 22 60.0 90.0 0
3 1066528 33 90.0 120.0 0
4 1066528 44 120.0 150.0 0

time_since_start_hours hr_mean hr_std hr_iqr hr_min ... \
0 0.000000 52.166667 0.897527 0.0 51.0 ...
1 0.008333 50.833333 1.950783 4.0 49.0 ...
2 0.016667 52.400000 1.356466 1.0 50.0 ...
3 0.025000 53.833333 0.372678 0.0 53.0 ...
4 0.033333 60.500000 4.500000 8.0 56.0 ...

acc_mag_max acc_mag_diff_std hr_mean_roll15 hr_std_roll15 \
0 1.017645 0.002779 NaN NaN
1 1.004546 0.001883 52.166667 0.897527
2 1.005295 0.001951 51.500000 1.424155
3 1.008659 0.002374 51.800000 1.401592
4 1.006992 0.002160 52.308333 1.144364

acc_mag_mean_roll15 acc_mag_std_roll15 hr_mean_roll120 hr_std_roll120 \
0 NaN NaN NaN NaN
1 1.000029 0.001947 52.166667 0.897527
2 1.000041 0.001725 51.500000 1.424155
3 1.000067 0.001666 51.800000 1.401592
4 1.000057 0.001751 52.308333 1.144364

acc_mag_mean_roll120 acc_mag_std_roll120
0 NaN NaN
1 1.000029 0.001947
2 1.000041 0.001725
3 1.000067 0.001666
4 1.000057 0.001751

```

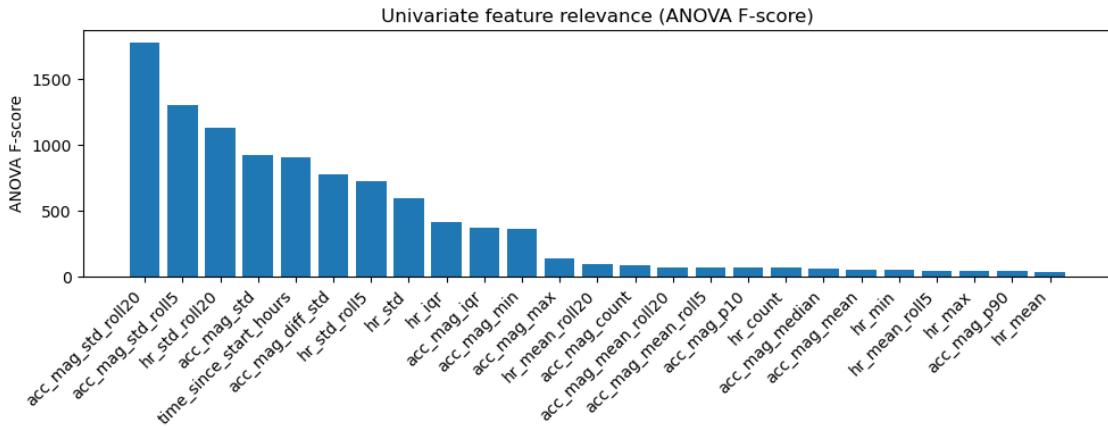
[5 rows x 30 columns]

4.5 Feature relevance and visualization

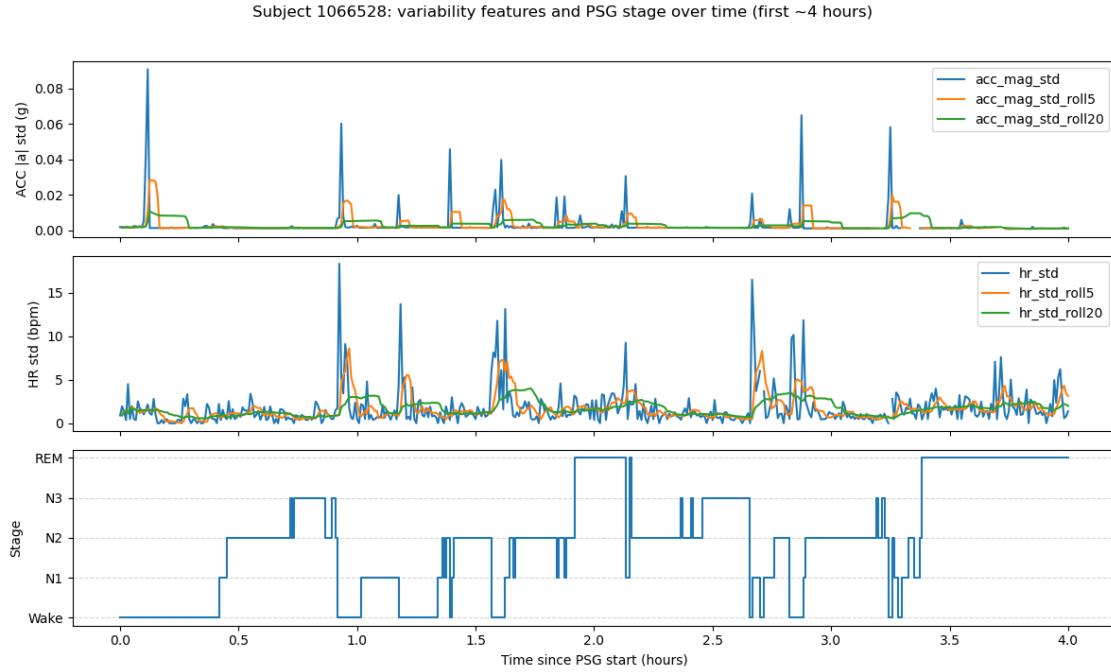
Before modeling, we do two quick sanity checks:

- 1) **Univariate feature relevance (ANOVA F-score)**: how strongly each feature differs across sleep stages (one feature at a time).
- 2) **Single-subject timeline**: how a couple of key features evolve over time alongside the PSG stage labels.

These visualizations are descriptive (not “proof” of separability) and help validate the pipeline before ML.



	feature	F_score	p_value
24	acc_mag_std_roll120	1781.548135	0.000000e+00
23	acc_mag_std_roll15	1300.868806	0.000000e+00
20	hr_std_roll120	1136.074761	0.000000e+00
7	acc_mag_std	926.236595	0.000000e+00
16	time_since_start_hours	909.576551	0.000000e+00
14	acc_mag_diff_std	781.981964	0.000000e+00
19	hr_std_roll15	727.759319	0.000000e+00
1	hr_std	595.518177	0.000000e+00
2	hr_iqr	420.137748	0.000000e+00
9	acc_mag_iqr	375.699593	2.950669e-315
12	acc_mag_min	366.777206	6.272101e-308
13	acc_mag_max	144.082104	4.462149e-122
18	hr_mean_roll120	98.912676	1.034501e-83
15	acc_mag_count	93.162596	8.167400e-79
22	acc_mag_mean_roll120	76.561146	1.162256e-64
21	acc_mag_mean_roll15	75.370355	1.205480e-63
10	acc_mag_p10	73.628140	3.694619e-62
5	hr_count	73.305709	6.961063e-62
8	acc_mag_median	59.866863	2.053812e-50
6	acc_mag_mean	57.542436	1.980793e-48
3	hr_min	52.692967	2.732684e-44
17	hr_mean_roll15	50.168934	3.899823e-42
4	hr_max	50.130534	4.205529e-42
11	acc_mag_p90	42.451121	1.502108e-35
0	hr_mean	41.436253	1.102355e-34



5. Modeling

We evaluate two label granularities using **subject-wise 5-fold GroupKFold** (leakage-aware sleep staging using the PSG-aligned epoch table):

Target

- **5-class:** Wake / N1 / N2 / N3 / REM
- **3-class:** Wake / NREM (N1+N2+N3) / REM

Features per-epoch summaries + causal history features (past-only rolling means).

Validation subject-wise 5-fold GroupKFold (group = `subject_id`) to avoid subject leakage.

Models (no tuning):

- **Logistic Regression (L2, class_weight="balanced")** (baseline)
- **HistGradientBoostingClassifier** (strong classical baseline)

We also inspect the **label distribution** to understand class imbalance.

```
5-class | CV: LR (L2, balanced): 100%|
          | 5/5
[00:03<00:00,  1.37it/s]
5-class | CV: HGB: 100%|
          | 5/5
[00:06<00:00,  1.24s/it]
5-class | Test-fold label distribution (% within each test fold):
```

	fold	Wake	N1	N2	N3	REM	n_test_epochs	n_test_subjects
0	0	9.82	6.70	47.28	13.97	22.22	5854	7
1	1	8.98	6.72	46.31	14.37	23.62	5254	6
2	2	8.23	6.59	48.17	14.06	22.94	5248	6
3	3	6.70	8.98	52.37	8.43	23.52	5043	6
4	4	12.20	5.52	51.49	11.82	18.97	5018	6

5-class | CV summary (mean ± std over folds, in %):

	run	model	macro_f1_mean	macro_f1_std	bal_acc_mean	\
1	5-class	LR (L2, balanced)	44.19	4.64	50.96	
0	5-class	HGB	43.62	4.54	42.40	
		bal_acc_std				
1		5.59				
0		4.76				

3-class | CV: LR (L2, balanced): 100% |

| 5/5

[00:00<00:00, 5.75it/s]

3-class | CV: HGB: 100% |

| 5/5

[00:03<00:00, 1.26it/s]

3-class | Test-fold label distribution (% within each test fold):

	fold	Wake	NREM	REM	n_test_epochs	n_test_subjects
0	0	9.82	67.95	22.22	5854	7
1	1	8.98	67.40	23.62	5254	6
2	2	8.23	68.83	22.94	5248	6
3	3	6.70	69.78	23.52	5043	6
4	4	12.20	68.83	18.97	5018	6

3-class | CV summary (mean ± std over folds, in %):

	run	model	macro_f1_mean	macro_f1_std	bal_acc_mean	\
0	3-class	HGB	55.71	4.60	54.12	
1	3-class	LR (L2, balanced)	54.73	4.89	61.06	
		bal_acc_std				
0		4.37				
1		6.75				

6. Results & reporting

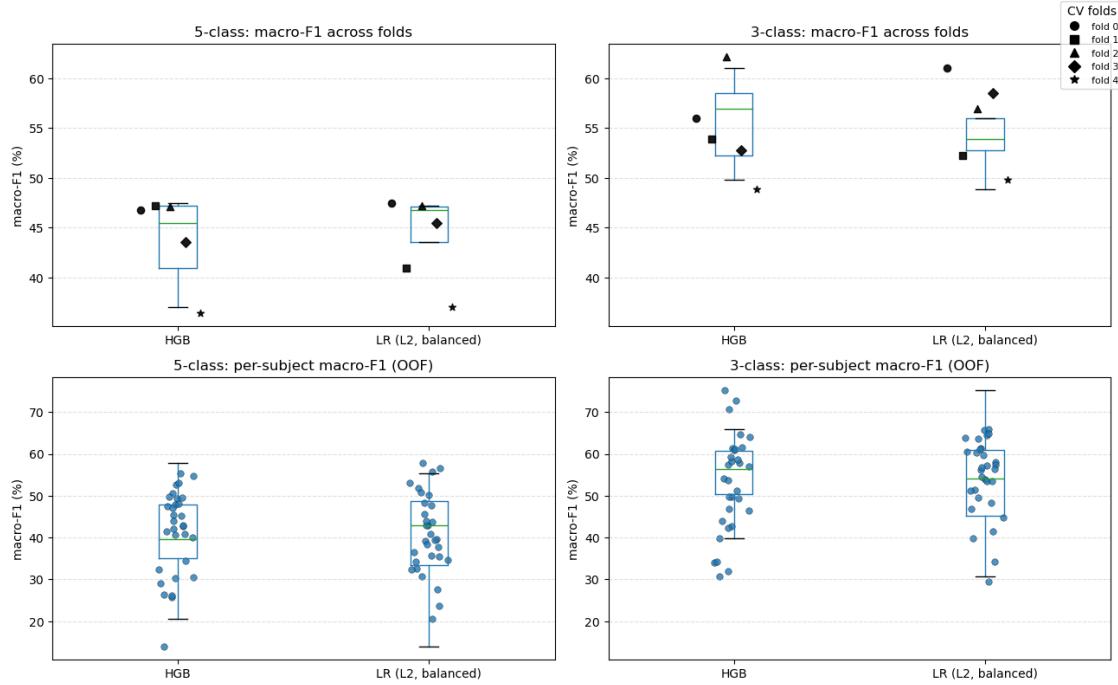
This section analyzes the **out-of-fold (OOF)** predictions produced by the subject-wise CV loop.

We report results for: - **5-class staging** (Wake / N1 / N2 / N3 / REM) - **3-class staging** (Wake / NREM / REM)

6.1 Performance summary

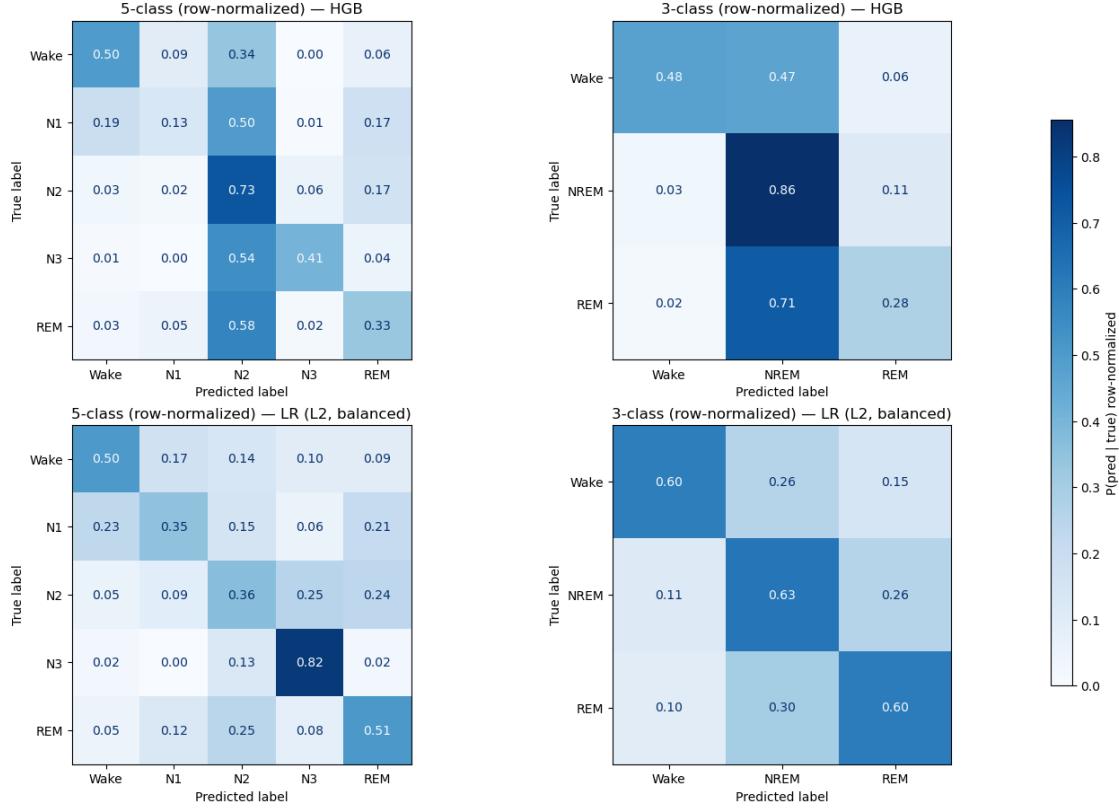
We summarize performance as:

- **Across folds:** macro-F1 variability across the 5 subject-wise folds.
- **Across subjects:** per-subject macro-F1 computed from out-of-fold predictions.



- **3-class (Wake/NREM/REM) is substantially easier than 5-class**, with macro-F1 shifting up for both models.
- **HGB is slightly better and more consistent across folds**, especially for 3-class; LR remains competitive.
- **Fold-to-fold variance is moderate** (a few harder folds appear as outliers), suggesting reasonable stability under subject-wise CV.
- **Per-subject variance is large** in both settings: some subjects are consistently much harder/easier, which is typical in wearable sleep due to signal quality and individual differences.
- **Takeaway:** for a wearable-style pipeline, **3-class staging** is a strong, realistic target; **5-class staging** remains challenging with lightweight wrist ACC/HR features.

6.2 Confusion matrices



- **5-class is difficult:** most errors happen between neighboring sleep stages. In particular, **N1** is rarely recognized cleanly and is often confused with **N2** (and sometimes Wake/REM).
- **N2 dominates the 5-class predictions:** both models tend to “collapse” uncertain epochs into **N2**, and **N3/REM** are frequently pulled toward N2 as well (more so for HGB here).
- **3-class is much clearer:** the main diagonal is stronger for both models, confirming that wrist signals separate **Wake vs NREM vs REM** better than fine-grained staging.
- **Key remaining confusion in 3-class: REM vs NREM** is the hardest boundary (especially for HGB), while LR appears **more balanced** between NREM and REM here.
- **Overall:** these matrices explain why **macro-F1 improves a lot** when moving from 5-class to 3-class: fewer ambiguous boundaries and less reliance on N1/N2/N3 separation.